

How Much Tensor Structure Should a Transformer FFN Have?

Routed Block-CP Feedforward Networks Between SwiGLU and Dense Tucker

Aniket Deshpande
University of Illinois Urbana-Champaign
aniket4@illinois.edu

June 2026

Abstract

Transformer feedforward networks hold most of a language model’s parameters, but their multiplicative interactions lack the structural vocabulary that makes attention analyzable. We start from an exact observation: a SwiGLU block is a *routed CP tensor model* — each hidden channel contributes a fixed rank-one interaction atom $u_j \otimes w_j \otimes g_j$, routed by an input-dependent sigmoid coefficient. The superdiagonal latent core this implies is both what atomizes the layer and what restricts every route to a single rank-one interaction. Rather than replacing it with an arbitrary dense Tucker core, we study the family that interpolates between the two: LL1 (block-term) feedforward networks, in which one gate routes a rank- L block of interactions, so that the per-route rank — exactly the quantity priced by an aligned-width separation theorem — becomes a hyperparameter, and, at fixed parameters, routing diversity is traded for interaction atoms. In controlled teacher–student recovery, routes and per-route rank are independently binding: students recover a routed tensor map to machine precision exactly when both suffice, and the nominally most expressive dense core loses every matched-budget comparison. On real pretrained FFN layers (Qwen2.5-0.5B), LL1 students with $L \approx 4$ –8 compress the layer’s function better than rank-one atoms in every layer and budget tested, and far better than a dense core. Trained from scratch at matched parameters and FLOPs (52.5M parameters, 100M tokens, 3 seeds), SwiGLU, LL1($L=4$), and dense Tucker tie on validation loss — while LL1 trains 5–7% faster than SwiGLU and $2.1\times$ faster than Tucker, and an unconstrained Tucker core spontaneously learns per-route stable rank ≈ 4 , the same value at which LL1’s blocks saturate and the distillation optimum sits. Interpretability proxies are a deliberate negative: no architecture routes sparsely, and LL1’s learned blocks recur across seeds no better than chance. Our results suggest that the right amount of tensor structure in a transformer FFN is a measurable quantity — about rank four per route at this scale — and that structured block-CP parameterizations deliver it at zero loss cost, but that architectural structure alone does not purchase mechanistic interpretability.

1 Introduction

Attention has accumulated a rich structural vocabulary — heads, QK/OV circuits, induction patterns (Elhage et al., 2021; Olsson et al., 2022) — that makes it the best-understood component of transformers. The feedforward network, which holds the majority of parameters in modern architectures (Vaswani et al., 2017; Shazeer, 2020), has no comparable account: its multiplicative gating creates feature interactions that are hard to trace through weights, and most interpretability progress on MLPs has come from analyzing activations rather than the architecture itself (Bricken et al., 2023; Dunefsky et al., 2024).

The starting point of this paper is that for SwiGLU — the dominant FFN variant — an exact architectural account already exists, hidden in plain sight. Because $\text{SiLU}(z) = z \sigma(z)$, each SwiGLU channel is a rank-one bilinear interaction between two learned directions, scaled by a sigmoid gate: the layer is, exactly and pointwise in its input, a CP decomposition of a third-order interaction tensor with fixed factors and input-dependent routing weights (§2.1). Each hidden index contributes one *atom* (what interaction is available) and one *route* (when it is used). This is the gated counterpart of the bilinear-MLP observation of Pearce et al. (2025), who showed that GLUs *without* the element-wise nonlinearity are third-order tensors amenable to weight-based analysis.

The CP structure cuts both ways. Read as a Tucker decomposition, SwiGLU’s latent core is superdiagonal: the j -th main feature interacts only with the j -th gate. This restriction is what makes the decomposition atomized and essentially unique (Kruskal, 1977) — “CPD gives a sparse core for free” — but it also means one route controls exactly one rank-one interaction. An earlier draft of this project proposed the obvious mathematical relaxation, a dense learned core (a Tucker-FFN), and proved a separation theorem: simulating a Tucker block with aligned SwiGLU units costs $\sum_j \text{rank } V_j$ units, where V_j is the interaction matrix routed by gate j — generically k^2 versus $O(k)$ latent dimensions. We now believe the dense core is the wrong conclusion to draw from that theorem. A dense core is gauge-non-identifiable (its latent directions carry no intrinsic meaning), concentrates $\sim 90\%$ of the block’s parameters in interactions, and — as our experiments show — is dominated in practice. The right reading of the theorem is that *the per-route rank $\text{rank } V_j$ is the priced quantity*, and the constructive half of its proof is itself an architecture: realize each V_j as L rank-one atoms sharing one gate.

That architecture is the LL1 (multilinear rank- $(L, L, 1)$ block-term) decomposition (De Lathauwer, 2008) applied to the FFN’s interaction tensor, with the gate mode as the rank-one mode (§2.3). It nests SwiGLU exactly at $L=1$ and block-sparse Tucker in general, and it makes the comparison crisp at a fixed parameter budget: an LL1 block buys $\frac{3L}{2L+1} \times$ more interaction atoms than SwiGLU at the price of $\frac{3}{2L+1} \times$ fewer independent routes. *How much rank should a route control?* becomes an empirical question with a dial.

We answer it with four experiments at strictly matched parameters and FLOPs: synthetic teacher–student recovery where the true structure is known (§3.1); compression of real pretrained FFN layers (§3.2); from-scratch language-model training (§3.3) with measured throughput; and interpretability proxies plus an induction-circuit pilot (§3.4–3.5).

Our contributions:

1. **Framing (exact):** SwiGLU as a routed CP tensor field; the LL1 FFN as the rank-controlled family between routed CP and dense Tucker, nesting both; the route/atom budget identity; and the aligned-width theorem re-read as pricing per-route rank (Appendix A).
2. **Recovery (synthetic):** at fixed budget, routes and per-route rank bind independently — error is minimized exactly at the teacher’s block rank, with machine-precision recovery if and only if structure matches, and the dense core loses every cross-class comparison.
3. **Real FFN maps:** distilling Qwen2.5-0.5B layers, LL1 with $L \approx 4\text{--}8$ beats rank-one atoms in all nine layer $\times$ budget cells ($\sim 30\times$ seed noise) and the dense core loses by 35–130%.
4. **End-to-end:** at 52.5M parameters / 100M tokens (3 seeds), LL1($L=4$), SwiGLU, and Tucker tie on loss; LL1 is 5–7% faster than SwiGLU and $2.1\times$ faster than Tucker; and the unconstrained Tucker core spontaneously learns per-route stable rank ≈ 4 — the value LL1 hard-codes.

5. **Interpretability (negative):** measured proxies — routing diffuseness, top- k decomposability, ablation locality, cross-seed factor recurrence — do not support the hypothesis that block structure improves mechanistic accessibility; LL1’s blocks recur across seeds at chance. The induction pilot shows no effect of FFN structure on circuit emergence, with a single-seed observation of an alternative dense-core copying circuit.

2 The routed tensor view of gated FFNs

Setup. At a fixed token position, let $x \in \mathbb{R}^d$ be the residual-stream vector entering the feedforward block. The SwiGLU FFN (Shazeer, 2020) computes

$$y = U^\top h, \quad h = (W^\top x) \odot \text{SiLU}(G^\top x), \quad (1)$$

with $W, G \in \mathbb{R}^{d \times m}$ and $U \in \mathbb{R}^{m \times d}$. Writing w_j, g_j for the j -th columns of W, G and u_j for the j -th column of $U^\top$, the j -th hidden coordinate is $h_j(x) = (w_j^\top x) \text{SiLU}(g_j^\top x)$.

2.1 SwiGLU is an exactly routed CP tensor field

Each SwiGLU channel multiplies two scalar projections of the residual stream. The SiLU nonlinearity factorizes exactly, $\text{SiLU}(z) = z \sigma(z)$ with $\sigma(z) = (1 + e^{-z})^{-1}$, so

$$h_j(x) = \underbrace{(w_j^\top x)(g_j^\top x)}_{\text{rank-one bilinear interaction}} \cdot \underbrace{\sigma(g_j^\top x)}_{\alpha_j(x) \in (0,1)}. \quad (2)$$

Defining the input-dependent interaction tensor $A(x) \in \mathbb{R}^{d \times d \times d}$ by $y_o = \sum_{a,b} A_{oab}(x) x_a x_b$, substituting (2) into (1) gives

$$A(x) = \sum_{j=1}^m \alpha_j(x) u_j \otimes w_j \otimes g_j, \quad \alpha_j(x) = \sigma(g_j^\top x). \quad (3)$$

This is a canonical polyadic (CP) decomposition (Kolda and Bader, 2009) of $A(x)$ with *shared, input-independent factors* and *input-dependent routing weights* $\alpha_j(x)$. The decomposition is exact and holds pointwise in x ; no approximation is involved. We call $u_j \otimes w_j \otimes g_j$ the j -th *interaction atom* (what interaction is available) and $\alpha_j(x)$ its *route* (when it is used). The bilinear $\alpha \equiv 1$ limit of this picture is the bilinear MLP analyzed by Pearce et al. (2025); the routed form keeps the gate inside the structural description rather than discarding it.

Viewed as a Tucker decomposition with latent dimensions $r = s = m$ and factor matrices $(P, Q, R) = (W, G, U^\top)$, the CP form (3) corresponds to a *superdiagonal core*: $C_{oij} = \delta_{oi} \delta_{ij}$. This restriction cuts both ways. It atomizes the layer — every hidden index is one (atom, route) pair, and CP decompositions are essentially unique under mild conditions (Kruskal, 1977), so the atoms are identifiable objects. But it also means one route controls exactly one rank-one interaction: feature w_i never interacts with gate g_j for $i \neq j$.

2.2 Dense Tucker and the per-route interaction matrix

The unconstrained relaxation drops the superdiagonal restriction. A *Tucker-core FFN* with latent projections $P, Q \in \mathbb{R}^{d \times r}$, output basis $R \in \mathbb{R}^{d \times s}$, and learned core $C \in \mathbb{R}^{s \times r \times r}$ computes $p = P^\top x$, $q = Q^\top x$,

$$z_o = \sum_{i,j=1}^r C_{oij} p_i \text{SiLU}(q_j), \quad y = Rz. \quad (4)$$

Grouping by gate index exposes the structure this paper is about:

$$y(x) = \sum_{j=1}^r V_j p \text{ SiLU}(q_j), \quad V_j := R C^{(j)} \in \mathbb{R}^{d \times r}, \quad C_{oi}^{(j)} = C_{oij}. \quad (5)$$

Each latent gate q_j routes a *matrix* V_j of output-by-main interactions. SwiGLU is the corner case $\text{rank } V_j = 1$ for all j ; a generic Tucker core has $\text{rank } V_j = \min(r, s)$.

Two properties make the dense core suspicious as an interpretability target. First, *gauge freedom*: for any invertible $M \in \mathbb{R}^{r \times r}$, replacing $P \mapsto PM^{-\top}$ and $C^{(j)} \mapsto C^{(j)}M$ for all j leaves the function unchanged, so individual latent directions p_i and individual core entries carry no intrinsic meaning — only the gate directions $Q_{:j}$ and the subspaces $\text{range}(V_j)$ are well defined. Second, the core costs sr^2 parameters; at the configurations we study it absorbs over 90% of the block’s budget.

The aligned-width theorem, reinterpreted. Prior work on this codebase proved an exact separation: if $[P \ Q]$ has full column rank, an *aligned* SwiGLU (units restricted to main vectors in $\text{span}(P)$ and gate vectors from the columns of Q) that exactly realizes (5) must have width $m \geq \sum_j \text{rank } V_j$, and a rank decomposition of each V_j attains the bound (Appendix A). Generically this is $r \min(r, s)$, i.e. k^2 in the balanced case — the result that originally motivated dense Tucker as “ $k \times$ cheaper than SwiGLU.” The reading we adopt here is different: *the theorem says the operative variable is the per-route rank $\text{rank } V_j$, not the presence of a dense core.* The constructive half of the proof — realize each V_j as a sum of $\text{rank } V_j$ rank-one terms sharing the same gate — is itself an architecture. That architecture is the subject of this paper.

2.3 LL1/block-CP: one route, one rank- L block

A third-order tensor admits a *decomposition in multilinear rank- $(L, L, 1)$ terms* (LL1; a structured block-term decomposition) if it can be written $T = \sum_{b=1}^B (A_b B_b^\top) \otimes c_b$: each term is rank- L in two modes and rank-one in the third (De Lathauwer, 2008; Vervliet et al., 2016). Taking the *gate mode* as the rank-one mode yields the FFN we study. The **LL1 FFN** with B blocks of rank L has, per block, a gate direction $g_b \in \mathbb{R}^d$, a main factor $A_b \in \mathbb{R}^{d \times L}$, and an output factor $U_b \in \mathbb{R}^{d \times L}$:

$$y(x) = \sum_{b=1}^B U_b (A_b^\top x) \text{ SiLU}(g_b^\top x). \quad (6)$$

Expanding into atoms, $y(x) = \sum_b \sum_{l=1}^L u_{b,l} (a_{b,l}^\top x) \text{ SiLU}(g_b^\top x)$: *grouped CP*, L rank-one atoms sharing one route. The interaction tensor is the routed LL1 tensor $A(x) = \sum_b \alpha_b(x) (U_b A_b^\top) \otimes g_b$, and the per-route matrix is $V_b = U_b A_b^\top$ with $\text{rank } V_b \leq L$ *by construction*. The aligned-width theorem’s control variable has become a hyperparameter.

The family nests both endpoints exactly. At $L = 1$, $B = m$, (6) is SwiGLU (verified to machine precision in our implementation); a single block with $L = r$ is a (one-gate) dense slice; and an LL1 FFN equals a Tucker FFN whose core is block-superdiagonal. LL1 decompositions also retain essential-uniqueness guarantees under mild conditions (De Lathauwer, 2008, 2011), unlike the gauge-free dense core.

What the parameter budget buys. The comparison that matters is at matched parameters:

$$N_{\text{SwiGLU}} = 3dm, \quad N_{\text{LL1}} = dB(2L + 1), \quad N_{\text{Tucker}} = d(2r + s) + sr^2. \quad (7)$$




Figure 1: The architecture family, indexed by per-route rank. **A)** A SwiGLU channel: two projections of the residual stream are multiplied and scaled by the route $\alpha_j(x) = \sigma(g_j^\top x)$; the channel contributes the rank-one atom $u_j \otimes w_j \otimes g_j$. **B)** An LL1 block: one route scales a rank- L bundle $V_b = U_b A_b^\top$ of interactions (L atoms sharing a gate). **C)** A dense Tucker core couples every latent main feature to every latent gate; each of its r routes carries a generically full-rank V_j , and the latent basis is not identifiable. **D)** At a fixed parameter budget the family trades routes against per-route rank (rectangles drawn to scale for the LM configuration of Table 1; area $\propto$ interaction atoms).

architecture	config	routes	per-route rank	atoms ($\sum_j \text{rank } V_j$)	FFN params/layer
SwiGLU	$m=1493$	1493	1	1493	2,293,248
LL1 $L=2$	$B=896$	896	2	1792	2,293,760
LL1 $L=4$	$B=498$	498	4	1992	2,294,784
LL1 $L=8$	$B=263$	263	8	2104	2,289,152
LL1 $L=16$	$B=136$	136	16	2176	2,297,856
dense Tucker	$r=s=128$	128	128	16384*	2,293,760

Table 1: Matched-budget configurations at $d = 512$ (the LM setting of §3.3). LL1 converts routing diversity into interaction atoms at fixed budget. *Tucker’s atom count is the aligned-SwiGLU width its generic core would require (Theorem 1); its parameters are concentrated in the core ($sr^2 = 2.10M$ of 2.29M).

At a fixed budget N , SwiGLU affords $m = N/3d$ atoms, each with a private route; LL1 affords $BL = \frac{3L}{2L+1} \cdot \frac{N}{3d}$ atoms — up to $1.5\times$ more — but only $B = \frac{3}{2L+1} \cdot \frac{N}{3d}$ routes. *LL1 trades routing diversity for interaction atoms.* An equivalent statement: an $LL1(B, L)$ block is exactly a width- BL SwiGLU whose gate vectors are tied in groups of L . Which side of the trade wins is an empirical question about what trained FFNs actually need, and it is the question our experiments are designed to answer. A prior observation sharpens the stakes: in a dense Tucker LM trained from scratch, the learned per-gate stable rank $\|V_j\|_F^2 / \|V_j\|_{\text{op}}^2$ concentrates around 4 — far below the generic value r , far above the SwiGLU value 1. The trained dense core already behaves like an LL1 model; LL1 makes that behavior an explicit and cheaper parameterization (with essential-uniqueness guarantees at the decomposition level, though §3.4 shows these do not translate into cross-seed stability of learned factors).

3 Experiments

All architectures are compared at matched parameter and FLOP budgets (Table 1; budgets match to $\pm 0.2\%$), with identical optimizers, schedules, and data. Code, configurations, and per-seed results accompany the paper; full hyperparameters are in Appendix B.




Figure 2: Synthetic teacher-student recovery at a fixed 9,216-parameter student budget ($d=64$; 3 seeds; line = best seed, band = seed range; log-log). Left: a CP teacher (32 routes) is recovered by rank-one students and degrades monotonically with student block rank L as gate-tying starves routes. Middle: an LL1 teacher (16 blocks, $L^*=4$) is recovered to machine precision exactly at $L=4$ (dash-dot vertical line) and fails in both directions — too few atoms below, too few routes above. Right: a generic dense-Tucker teacher is recovered only by the Tucker student (red dotted line at 10^{-10}). Matched structure wins everywhere.

3.1 Synthetic teacher-student recovery: structure must match, in both directions

Question. Is per-route rank a real capacity dial — does a student recover a routed tensor map if and only if its routes and per-route ranks suffice?

Setup. Teachers at $d = 64$ with unit-variance outputs: a CP teacher (32 rank-one routed atoms), an LL1 teacher ($B^*=16$ blocks of rank $L^*=4$, i.e. 64 atoms on 16 routes), and a dense Tucker teacher ($r=s=16$, generic full-rank slices). Students of every family are trained on Gaussian inputs by MSE at a fixed 9,216-parameter budget, 3 seeds; we report validation MSE relative to output variance (best of 3, capacity reading; bands span seeds).

Result. Figure 2 (left to right): on the CP teacher, $L=1$ students recover it (rel. MSE $\sim 10^{-4}$) and error rises monotonically with L (gate-tying starves routes: an $L=16$ student has only $B=4$ routes for a 32-route teacher). On the LL1 teacher the error minimum sits exactly at $L = L^* = 4$ (machine precision, 2×10^{-13} , 2/3 seeds), with failure on both sides: $L < 4$ students lack atoms ($48, 58 < 64 = \sum_b \text{rank } V_b$), the $L=8$ student lacks routes ($8 < 16$). On the Tucker teacher only the Tucker student succeeds (1.1×10^{-10}); every CP/LL1 student plateaus at rel. MSE ≥ 0.2 .

Interpretation. Routes and per-route rank are *independently binding* resources at fixed budget, exactly as the counting argument predicts, and each family attains machine-precision recovery of its own class — the representations are optimization-reachable, not just existence proofs. Expressivity-as-superset does not survive a parameter budget: the Tucker student, which contains all other classes, loses on every non-Tucker teacher because its core consumes the budget that CP/LL1 students spend on routes.

Caveat. Gaussian inputs, structure-matched teachers, $d = 64$. This experiment calibrates the dial; it says nothing about which structure real data prefers.

3.2 Compressing real FFN maps: pretrained layers prefer small $L > 1$

Question. Which tensor structure best describes the input-output function a *real*, *pretrained* SwiGLU FFN computes?

Setup. We capture FFN inputs and outputs for layers $\{4, 12, 20\}$ of Qwen2.5-0.5B ($d=896$, $m=4864$; 13.07M parameters per FFN) on $\approx 120K$ tokens of FineWeb-Edu, and fit students of every family at compression budgets $\{0.6, 1.2, 2.4\}M$ parameters (4.6–18.4% of the teacher) by




Figure 3: Distilling pretrained Qwen2.5-0.5B FFN layers into matched-budget students (relative validation MSE at the 1.2M-parameter budget; 2 seeds, spread ± 0.001 ; log-log). In every layer the LL1 student improves on the SwiGLU/CP student (dashed blue) with a shallow optimum at block rank $L \approx 4-16$, and the dense Tucker student (dotted red) is far worse. Real FFN maps contain routed low-rank block structure.

MSE. 2 seeds; the seed spread is ± 0.001 rel. MSE throughout, so all reported gaps are $\gtrsim 30\times$ noise.

Result. The ordering $\text{LL1}(L \approx 4-16) < \text{LL1}(2) < \text{CP} \ll \text{dense Tucker}$ holds in all nine (layer, budget) cells (Figure 3). LL1’s gain over CP is modest but systematic (2–4% rel. MSE; e.g. layer 12 at 2.4M: 0.348 vs 0.361), with a shallow optimum at $L \approx 4-8$. The $L=1$ control matches SwiGLU to ± 0.001 in every cell. Dense Tucker is worse by 35–130% rel. MSE and does not improve — on layer 20 it *worsens* — with budget.

Interpretation. The function computed by a trained SwiGLU layer contains cross-atom structure that rank-one routed atoms capture inefficiently: at matched budget, spending parameters on rank-4–8 routed blocks (and accepting 3–10 $\times$ fewer routes) is consistently the better trade. The dense core is the wrong basis for this function class at compression budgets.

Caveat. This is the compression regime; it measures the structure of the teacher’s function, not end-to-end trainability (next subsection). One model family (Qwen2.5) and one dataset.

3.3 Matched-budget language modeling: a three-way tie on loss; a $2\times$ split on speed

Question. Do these representational differences survive end-to-end training on language?

Setup. LLaMA-style decoders (52.5M parameters, $d=512$, 8 layers) trained from scratch on 100M FineWeb-Edu tokens with identical hyperparameters, differing only in FFN (Table 1); 3 seeds for SwiGLU, LL1($L=4$), and Tucker; 1 seed for the remaining L values. Tucker uses its best-known (diagonal warm-start) initialization. Throughput is measured separately on an idle A40.

Result. Final validation loss: LL1($L=4$) 4.7472 ± 0.0041 , SwiGLU 4.7542 ± 0.0104 , Tucker 4.7626 ± 0.0071 (mean $\pm$ std over 3 seeds; Figure 4). LL1 vs. SwiGLU is a statistical tie (Welch $t \approx 1.1$); LL1 vs. Tucker separates ($t \approx 3.1$). The L -sweep is flat with a shallow optimum away from the corners (single seeds: $L=1$: 4.765 — matching SwiGLU’s range and confirming the $L=1 \equiv \text{SwiGLU}$ control end-to-end; $L=2$: 4.740; $L=8$: 4.742; $L=16$: 4.750). Training throughput at matched FLOPs (4.59×10^6 MACs/token/layer for every architecture): LL1 75.3K tokens/s, SwiGLU 71.4K, Tucker 36.5K — LL1 is 5–7% faster than SwiGLU (its gate projection is a factor $\frac{2L+1}{3}$ narrower, and the gating multiply touches B rather than m columns) and $2.06\times$ faster than Tucker, whose core contraction is GEMM-unfriendly.

Interpretation. At this scale the route/atom trade is loss-neutral on language modeling: tripling per-route rank while tripling down on routes neither helps nor hurts. The differences that




Figure 4: Matched-budget language modeling (52.5M parameters, 100M FineWeb-Edu tokens). Left: final validation loss across the LL1 block-rank sweep (error bar = std over 3 seeds at $L=4$; other L single-seed), with SwiGLU (blue) and dense Tucker (red) 3-seed bands. Every $L \geq 2$ sits at or below the SwiGLU mean; the sweep is flat. Right: validation-loss curves (mean $\pm$ std over 3 seeds) — the three families are indistinguishable through training.

survive are operational: LL1 buys a throughput edge and a structurally bounded per-route rank at zero loss cost, while the dense core costs $2\times$ wall-clock for nominally identical FLOPs and slightly worse loss.

Caveat. 52.5M parameters, 100M tokens, one dataset, hyperparameters inherited from SwiGLU tuning; the L -sweep is single-seed away from $L=4$.

3.4 Interpretability proxies: the computation has rank ≈ 4 per route, and no architecture routes sparsely

Question. What does block rank buy or cost in measurable structure?

Setup. On the trained LMs: per-token unit contributions and their exp-entropy (“effective active units”), all-layer top- k decomposition curves, single-unit ablations, weight-based per-gate stable ranks, and cross-seed factor matching with shuffled-weight nulls (Appendix B).

Result. (i) The trained dense Tucker models use per-gate stable rank 3.97/3.99/3.97 (three seeds; max 5.3) of an available 128; LL1’s realized rank tracks but under-saturates its cap (1.83 at $L=2$, 3.26 at $L=4$, 5.62 at $L=8$). With the distillation knee at $L \approx 4-8$, three independent measurements give the same answer: *the natural per-route rank of FFN computation at this scale is $\approx 4-6$* . (ii) No architecture routes sparsely: effective active units per token are 49% of capacity for SwiGLU (733/1493), 63% for LL1 (312/498), 64% for Tucker (82/128); to stay within 0.03 nats of base loss, SwiGLU needs $\sim 34\%$ of its atoms per token, LL1 $\sim 51\%$ of blocks, Tucker $\sim 100\%$ of gates (Figure 5). (iii) Single-unit ablations are negligible for all (max $\Delta\text{loss} \leq 0.0016$ nats). (iv) Cross-seed recurrence: SwiGLU’s joint atoms match across seeds at $3.0\times$ the shuffled null (cosine 0.268 vs. 0.089); gauge-invariant per-route matrices recur weakly for SwiGLU ($1.7\times$ null) and Tucker ($1.6\times$) but are *at chance* for LL1; gate directions are at chance for every architecture.

Interpretation. The rank-4 convergence is the cleanest structural fact this sprint produced, and it argues for LL1 as the *honest parameterization*: it hard-codes the rank the unconstrained model learns anyway, at lower cost. But the interpretability proxies do not support the hypothesis that this buys mechanistic legibility: routing stays diffuse, blocks are individually dispensable, and LL1’s learned blocks recur across seeds no better than chance — the block partition apparently




Figure 5: Top- k decomposability of the trained LMs: excess validation loss when every layer’s FFN output is replaced by its k largest-contribution routed units per token (left: absolute k ; right: fraction of available units; mean over 3 seeds for SwiGLU/LL1, 2–3 for Tucker; log–log). SwiGLU is the most truncatable relative to capacity, Tucker the least; LL1 sits between, with $3\times$ fewer (larger) objects than SwiGLU at equal excess loss.

adds assignment freedom that atom-level CP does not have.

Caveat. Proxies measure routing statistics and weight structure, not semantic legibility; no human or automated interpretation of atoms/blocks was attempted. The Tucker stable-rank value is conditional on the diagonal warm-start initialization (the only one that reaches SwiGLU-level loss; prior work on this codebase found random-init Tucker converges to stable rank ≈ 27 at worse loss); the distillation knee and the LL1 cap-saturation pattern are initialization-independent corroborations.

3.5 Circuit pilot: induction emerges identically; one dense-core seed finds another circuit

Question. Does FFN tensor structure change the emergence or mechanism of a known attention circuit?

Setup. Repeated-sequence induction task ($[s; s]$, fresh random s each sample, vocab 128, $|s|=64$); 2-layer transformers ($d=128$) identical except the FFN (matched 131K-parameter budgets; attention-only control); 3 seeds; we track second-half accuracy, the canonical induction score (attention mass at offset $|s|-1$), attention-offset profiles, and head/FFN ablations.

Result. Emergence is architecture-independent: every FFN variant reaches 100% accuracy in 140–180 steps (attention-only: 80). All SwiGLU/LL1 runs and 2/3 Tucker runs converge to textbook induction attention (offset-63 mass 0.93–1.00). One Tucker seed reaches 100% accuracy with *no* canonical induction head: its layer-2 heads attend at diffuse offsets (mass ≤ 0.19 spread over offsets 37–56), ablating its best head changes nothing, and its solution runs through the FFN. A methodological caution: zeroing all FFNs collapses *every* architecture (including SwiGLU, to 13%), so FFN-bypass is not a valid circuit-attribution test for jointly trained models; the offset profiles are the discriminating measurement.

Interpretation. A negative result for the hypothesis that FFN tensor structure modulates circuit emergence at this scale, plus a single-seed existence proof that the dense-core FFN admits a qualitatively different copying implementation. We report the latter as an observation ($n=1$ of 3 seeds), not a claim about frequency.

Caveat. Tiny models, one task, three seeds; the alternative-circuit observation needs a basin-frequency study.

4 Discussion

What is established. Per-route interaction rank is a real capacity dial: at fixed parameters it binds in both directions in controlled recovery, and three independent measurements — the learned stable rank of an unconstrained Tucker core (≈ 3.98 , replicated), the saturated rank of capped LL1 blocks (3.26 of 4), and the compression optimum on pretrained FFN maps ($L \approx 4-8$) — agree that LM FFN computation at this scale uses rank ≈ 4 per route. Given that, the dense Tucker core is a dominated design point: it spends 91% of its budget parameterizing interactions it will not use, costs $2\times$ wall-clock at matched FLOPs, and ends nominally last on validation loss. The structured middle of the family is free: LL1 at $L=4$ matches SwiGLU’s loss with a small throughput edge while fixing, by construction, the rank structure the unconstrained model discovers by itself.

What is plausible but not established. That LL1’s advantages grow with scale (its compression edge on real maps suggests the function class is right, but 52.5M-parameter training cannot confirm it); that the route/atom trade has a scale- or data-dependent optimum away from $L=1-4$; that rank-4 blocks are semantically coherent objects.

What failed. Our pre-registered interpretability hypothesis — that LL1 would land between SwiGLU and Tucker on diffuseness while keeping CP-like factor stability — is unsupported. Routing is diffuse for every architecture (49–64% of capacity effectively active per token); and on the one proxy where architectures separate, cross-seed factor recurrence, LL1 is the only family *at chance* under a gauge-invariant test that SwiGLU and Tucker both pass weakly. Identifiability of a decomposition given a fixed function evidently does not imply stability of the learned function’s factors across training runs. The circuit pilot returned a clean null on emergence.

Limitations. One scale (52.5M / 100M tokens), one dataset and tokenizer, hyperparameters inherited from SwiGLU tuning (no per-architecture tuning; this likely flatters the baseline), Tucker evaluated with its best-known initialization and weight-decay exemption on the core (an asymmetry in its favor), single-seed L -sweep away from $L=4$, compression-regime distillation only, and statistical interpretability proxies rather than semantic evaluation of learned blocks.

Next step. The single most informative follow-up is a $10\times$ -scale repeat of the L -sweep (≥ 300 M parameters, ≥ 1 B tokens, per-architecture learning-rate tuning, ≥ 3 seeds) testing whether the rank-4 optimum, the loss tie, and the throughput edge persist; second, a trained sparse variant (L1 on block routes) targeting the routing diffuseness that post-hoc truncation cannot remove.

5 Related work

Gated FFNs and bilinear MLPs. GLU-family FFNs outperform ReLU/GELU MLPs at matched scale and are standard in modern LMs (Shazeer, 2020). Pearce et al. (2025) study bilinear MLPs — GLUs with the element-wise nonlinearity removed — which are exactly third-order tensors, and show that eigendecomposition of those tensors yields interpretable low-rank structure; Doms and Wilhelm (2024) make the architectural case. Our object differs in keeping the gate:

SwiGLU is exactly a *routed* CP tensor field whose routing is input-dependent, and prior experiments on this codebase show the routing is load-bearing (constant- α ablations degrade pretrained Qwen models by four to six orders of magnitude in perplexity), so the bilinear picture alone does not describe trained SwiGLU models. Jayakumar et al. (2020) situate gating within multiplicative interactions broadly; we study the specific tensor decompositions a gated FFN admits.

Tensor decompositions in neural networks. CP, Tucker, and block-term decompositions are classical (Kolda and Bader, 2009; De Lathauwer, 2008). In deep learning they are most often used for *compression* of trained weights (Novikov et al., 2015) or for expressivity analyses of idealized architectures (Cohen et al., 2016); factorization machines use low-rank multiplicative atoms for feature interactions (Rendle, 2010); Loconte et al. (2024) connect tensor factorizations to probabilistic circuits. Our use is different: the decomposition is not applied post hoc to a weight tensor — the FFN *is* the decomposition, and we ask which member of the CP $\rightarrow$ LL1 $\rightarrow$ Tucker family is the right parameterization to train.

Sparse dictionaries and weight-based interpretability. Sparse autoencoders and transcoders decompose activations into feature directions (Bricken et al., 2023; Dunefsky et al., 2024; Paulo et al., 2025). The routed-CP/LL1 view is complementary: atoms and routes are *architectural* objects fixed in the weights, with the input-dependence isolated in scalar routing coefficients — the same invariant-dictionary + input-dependent-coefficients factorization transcoders recover empirically.

Structured matrices for efficiency. Monarch and butterfly factorizations replace dense projections with products of block-diagonal or sparse factors (Dao et al., 2022, 2019); LoRA demonstrates that rank constraints are effective in transformers (Hu et al., 2022). These address the *factor matrices*; our axis is the *latent core* connecting them. The two are orthogonal and composable.

Circuit-level evaluation. Induction heads are the canonical recognizable circuit (Elhage et al., 2021; Olsson et al., 2022); we use a minimal induction task to ask whether FFN tensor structure changes the emergence or mechanism of an attention-side circuit.

References

- Trenton Bricken, Adly Templeton, Joshua Batson, Brian Chen, Adam Jermyn, Tom Conerly, et al. Towards monosemanticity: Decomposing language models with dictionary learning. *Transformer Circuits Thread*, 2023.
- Nadav Cohen, Or Sharir, and Amnon Shashua. On the expressive power of deep learning: A tensor analysis. In *Conference on Learning Theory*, 2016.
- Tri Dao, Albert Gu, Matthew Eichhorn, Atri Rudra, and Christopher Ré. Learning fast algorithms for linear transforms using butterfly factorizations. In *International Conference on Machine Learning*, 2019.
- Tri Dao, Beidi Chen, Nimit Sohoni, Arjun Desai, Michael Poli, Jessica Grogan, Alexander Liu, Aniruddh Rao, Atri Rudra, and Christopher Ré. Monarch: Expressive structured matrices for efficient and accurate training. In *International Conference on Machine Learning*, 2022.

- Lieven De Lathauwer. Decompositions of a higher-order tensor in block terms—part II: Definitions and uniqueness. *SIAM Journal on Matrix Analysis and Applications*, 30(3):1033–1066, 2008.
- Lieven De Lathauwer. Blind separation of exponential polynomials and the decomposition of a tensor in rank- $(l_r, l_r, 1)$ terms. *SIAM Journal on Matrix Analysis and Applications*, 32(4):1451–1474, 2011.
- Thomas Doods and Daniel Wilhelm. Weight-based decomposition: A case for bilinear MLPs. *arXiv preprint arXiv:2406.03947*, 2024.
- Jacob Dunefsky, Philippe Chlenski, and Neel Nanda. Transcoders find interpretable LLM feature circuits. *arXiv preprint arXiv:2406.11944*, 2024.
- Nelson Elhage, Neel Nanda, Catherine Olsson, Tom Henighan, Nicholas Joseph, Ben Mann, Amanda Askell, Yuntao Bai, Anna Chen, Tom Conerly, et al. A mathematical framework for transformer circuits. *Transformer Circuits Thread*, 2021.
- Edward J. Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yanzhi Li, Shean Wang, Lu Wang, and Weizhu Chen. LoRA: Low-rank adaptation of large language models. In *International Conference on Learning Representations*, 2022.
- Siddhant M. Jayakumar, Wojciech M. Czarnecki, Jacob Menick, Jonathan Schwarz, Jack Rae, Simon Osindero, Yee Whye Teh, Tim Harley, and Razvan Pascanu. Multiplicative interactions and where to find them. In *International Conference on Learning Representations*, 2020.
- Tamara G. Kolda and Brett W. Bader. Tensor decompositions and applications. *SIAM Review*, 51(3):455–500, 2009.
- Joseph B. Kruskal. Three-way arrays: Rank and uniqueness of trilinear decompositions, with application to arithmetic complexity and statistics. *Linear Algebra and its Applications*, 18(2): 95–138, 1977.
- Lorenzo Loconte, Antonio Mari, Gennaro Gala, Robert Peharz, Cassio de Campos, Erik Quaeghebeur, Gennaro Vessio, and Antonio Vergari. What is the relationship between tensor factorizations and circuits (and how can we exploit it)? *arXiv preprint arXiv:2409.07953*, 2024.
- Alexander Novikov, Dmitry Podoprikin, Anton Osokin, and Dmitry Vetrov. Tensorizing neural networks. In *Advances in Neural Information Processing Systems*, 2015.
- Catherine Olsson, Nelson Elhage, Neel Nanda, Nicholas Joseph, Nova DasSarma, Tom Henighan, Ben Mann, Amanda Askell, Yuntao Bai, Anna Chen, et al. In-context learning and induction heads. *Transformer Circuits Thread*, 2022.
- Gonalo Paulo, Stepan Shabalin, and Nora Belrose. Transcoders beat sparse autoencoders for interpretability. *arXiv preprint arXiv:2501.18823*, 2025.
- Michael T. Pearce, Thomas Doods, Alice Rigg, Jose M. Oramas, and Lee Sharkey. Bilinear MLPs enable weight-based mechanistic interpretability. In *International Conference on Learning Representations*, 2025.
- Steffen Rendle. Factorization machines. In *IEEE International Conference on Data Mining*, 2010.
- Noam Shazeer. GLU variants improve transformer. *arXiv preprint arXiv:2002.05202*, 2020.

Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N. Gomez, Łukasz Kaiser, and Illia Polosukhin. Attention is all you need. In *Advances in Neural Information Processing Systems*, 2017.

Nico Vervliet, Otto Debals, Laurent Sorber, Marc Van Barel, and Lieven De Lathauwer. Tensorlab 3.0. Available online at <https://www.tensorlab.net>, 2016.

An Yang et al. Qwen2.5 technical report. *arXiv preprint arXiv:2412.15115*, 2024.

A The aligned-width theorem

This appendix restates, with proof, the separation theorem from the earlier draft accompanying this codebase, and records its constructive half as the LL1 architecture. Throughout, $\phi = \text{SiLU}$; the only properties used are $\phi(0) = 0$ and $\phi(1) \neq 0$.

Lemma 1 (SwiGLU recovery). *Setting $r = s = m$, $P = W$, $Q = G$, $R = U^\top$, and $C_{oij} = \delta_{oi}\delta_{ij}$ in (4) reduces the Tucker-core FFN to the SwiGLU map (1).*

Proof. With the superdiagonal core, (4) collapses to $z_o = p_o \phi(q_o)$, so $z = (P^\top x) \odot \phi(Q^\top x)$, which is h from (1); then $y = Rz = U^\top h$. $\square$

Definition 1 (Aligned SwiGLU representation). *Given latent matrices $P, Q \in \mathbb{R}^{d \times r}$, an aligned width- m SwiGLU representation is a map*

$$\hat{y}(x) = \sum_{j=1}^r \sum_{\nu=1}^{k_j} u_{j\nu} (a_{j\nu}^\top p) \phi(q_j), \quad \sum_j k_j = m,$$

with $u_{j\nu} \in \mathbb{R}^d$, $a_{j\nu} \in \mathbb{R}^r$, $p = P^\top x$, $q = Q^\top x$. Each unit uses a main vector in $\text{span}(P)$ and one of the gate vectors $Q_{:j}$.

Theorem 1 (Diagonal bottleneck). *Assume $[P \ Q] \in \mathbb{R}^{d \times 2r}$ has full column rank. The minimum width of an aligned SwiGLU representation exactly realizing a Tucker-core FFN with per-gate coefficient matrices $\{V_j\}_{j=1}^r$ is $m_{\min} = \sum_{j=1}^r \text{rank } V_j$.*

Proof. Lower bound. By the rank assumption, $x \mapsto (p, q)$ is surjective onto $\mathbb{R}^r \times \mathbb{R}^r$, so an exact representation forces $\sum_j (V_j - \hat{V}_j) p \phi(q_j) = 0$ for all (p, q) , where $\hat{V}_j = \sum_\nu u_{j\nu} a_{j\nu}^\top$ has rank at most k_j . Fix j_0 and set $q_j = \delta_{jj_0}$; since $\phi(0) = 0$ and $\phi(1) \neq 0$, only the j_0 term survives and $V_{j_0} = \hat{V}_{j_0}$. Hence $k_j \geq \text{rank } V_j$ for every j , and $m = \sum_j k_j \geq \sum_j \text{rank } V_j$.

Tightness. Let $V_j = \sum_{\nu=1}^{\rho_j} u_{j\nu} a_{j\nu}^\top$ be any rank decomposition, $\rho_j = \text{rank } V_j$. Then $V_j p \phi(q_j) = \sum_\nu u_{j\nu} ((Pa_{j\nu})^\top x) \phi(Q_{:j}^\top x)$ is a sum of ρ_j aligned units sharing gate $Q_{:j}$. $\square$

Corollary 1 (Generic separation). *If R has full column rank and each slice $C^{(j)}$ is generic, then $\text{rank } V_j = \min(r, s)$ for all j and $m_{\min} = r \min(r, s)$; in the balanced case $r = s = k$, $m_{\min} = k^2$.*

Remark 1 (The constructive half is LL1). *The tightness construction — ρ_j rank-one units sharing the gate $Q_{:j}$ — is precisely one block of the LL1 FFN (6) with $L = \rho_j$. Theorem 1 therefore does not say “dense cores beat SwiGLU”; it says the cost of simulating a routed tensor map is the sum of its per-route ranks, and an architecture that fixes those ranks directly realizes the map at exactly that cost.*

B Experimental details

Hardware. All sprint experiments ran on two NVIDIA A40 GPUs (46 GB).

Synthetic teacher–student (§3.1). $d = 64$; teachers: CP = LL1($B=32, L=1$), LL1($B=16, L=4$), dense Tucker $r=s=16$ with slices resampled to full rank and core scaled by $1/r$; all factor entries i.i.d. $\mathcal{N}(0, 1/d)$; teacher outputs normalized to unit std. Students trained on 50,000 / validated on 8,000 inputs $x \sim \mathcal{N}(0, I_d)$; Adam, lr 3×10^{-3} cosine-annealed to 1%, batch 512, 5,000 steps, 3 seeds per cell. Budgets: L -sweep at $N=9216$ (within 8% for all archs; exact counts reported with results); budget sweep $\{0.25, 0.5, 1, 2\} \times N$.

Language modeling (§3.3). LLaMA-style decoder: $d=512$, 8 layers, 8 heads, RoPE, RMSNorm, tied embeddings, GPT-2 BPE (50,257), context 1024. FFN budget matched to SwiGLU $m=1493$ (2.293M/layer $\pm 0.2\%$): Tucker $r=s=128$; LL1 $(L, B) \in \{(1, 1493), (2, 896), (4, 498), (8, 263), (16, 136)\}$. AdamW $(\beta_1, \beta_2) = (0.9, 0.95)$, weight decay 0.1 (Tucker core 0.0, inherited from prior work), peak lr 3×10^{-4} , 200-step warmup, cosine to 10%, grad clip 1.0, batch 24×1024 tokens, bf16 autocast, 100M FineWeb-Edu (sample-10BT) tokens (≈ 4070 steps), validation on a held-out 128-sequence (131K-token) split with deterministic seed. Tucker uses the variance-preserving diagonal warm start ($C_{ooo}=1$, off-diagonal std $10^{-2}/r$), the only initialization that matched SwiGLU in prior work on this codebase. Seeds 0–2 for SwiGLU, Tucker, LL1($L=4$); seed 0 for the remaining L values.

Pretrained-layer distillation (§3.4). Teacher: Qwen2.5-0.5B (Yang et al., 2024) ($d=896$, $m=4864$, 13.07M params/FFN). We capture FFN input/output pairs on ≈ 120 K tokens of FineWeb-Edu for layers $\{4, 12, 20\}$, normalize outputs to unit std, and fit students at budgets $\{0.6, 1.2, 2.4\}$ M parameters (4.6–18.4% of the teacher layer) by MSE; Adam lr 2×10^{-3} , 4,000 steps, batch 512, 2 seeds; 90/10 train/val split.

Induction pilot (§3.5). Sequences $[s; s]$, s uniform over vocab 128, $|s| = 64$; 2-layer transformers, $d=128$, 4 heads; FFN budget 131,072 params (SwiGLU $m=341$; LL1 $L \in \{1, 4, 16\}$; Tucker $r=48$; attention-only control). AdamW lr 10^{-3} , weight decay 0.01, batch 128, 3,000 steps, 3 seeds. Induction score: mean attention mass from second-half queries at offset $|s| - 1$, max over layer-2 heads, on a fixed 256-sequence eval batch.

Interpretability proxies (§3.4). Contributions computed on 4096 randomly sampled token positions from a held-out 32-sequence validation set. Unit contributions: SwiGLU $c_j = |h_j| \|u_j\|$; LL1 $c_b = \|U_b r_b\| |\text{SiLU}(g_b^\top x)|$; Tucker $c_j = \|V_j p\| |\text{SiLU}(q_j)|$. Top- k curves replace every layer’s FFN output by the sum of its k largest-contribution units per token. Single-unit ablations zero one atom (SwiGLU), one block-gate (LL1), or one latent gate (Tucker; $\text{SiLU}(0) = 0$ removes the slice exactly) in one layer and re-evaluate full-model validation loss.

C Additional results




Figure 6: Synthetic recovery across student parameter budgets $\{0.25, 0.5, 1, 2\} \times 9216$ (mean over 3 seeds, band = seed range). The matched-structure advantage of Figure 2 persists at every budget; the dense Tucker student improves on non-Tucker teachers only marginally with budget because added parameters go into its core rather than into routes.




Figure 7: Induction pilot: second-half accuracy (left) and canonical induction score (right) versus training step (mean $\pm$ std over 3 seeds per architecture). All FFN variants learn the task at indistinguishable speed; the attention-only control is faster. The Tucker induction-score band is wide because one of its three seeds solves the task without the canonical attention pattern.




Figure 8: Layer-2 attention-offset profiles (mean attention mass from second-half query positions to each relative offset). Left: a representative SwiGLU run — all heads concentrate on the canonical induction offset $|s| - 1 = 63$ (mass 0.93–1.0). Right: the anomalous Tucker seed — attention is spread over non-canonical offsets (mass ≤ 0.19 per head), yet the model reaches 100% accuracy; ablating its best head changes nothing.